

# Predicting Diabetes Patient Readmission

Sheng Weng edit by mahdi akbari zarkesh

## Background

It was reported that in 2011 more than 3.3 million patients were readmitted in the US within 30 days of being associated with about \$41 billion in hospital costs. The need for readmission indicates that inadequate care time of first admission. The readmission rate has become an important metric measuring the overall quality

Diabetes is the 7th leading cause of death and affects about 23.6 million people in the US. 1.4 million Americans die every year. Hospital readmission being a major concern in diabetes care, over 250 million dollars was spent on patients in 2011. Early identification of patients facing a high risk of readmission can enable healthcare providers to conduct investigations and possibly prevent future readmissions.

In this project, I build a machine learning classifier model to predict diabetes patients with high risk of readmission (recall) is more desirable for hospitals because it is more crucial to correctly identify "high risk" patients who are at risk of readmission than identifying "low risk" patients.

## Dataset Description

The dataset represents 10 years (1999-2008) of clinical care at 130 US hospitals and integrated delivery networks representing 101766 diabetes patients and hospital outcomes. Information was extracted from the database based on the following criteria:

- It is an inpatient encounter (a hospital admission).
- It is a diabetic encounter, that is, one during which any kind of diabetes was entered to the system as a diagnosis.
- The length of stay was at least 1 day and at most 14 days.
- Laboratory tests were performed during the encounter.
- Medications were administered during the encounter.

The data contains such attributes as patient number, race, gender, age, admission type, time in hospital, medical history, physician, number of lab test performed, HbA1c test result, diagnosis, number of medication, diabetic medication, inpatient, and emergency visits in the year before the hospitalization, etc.

Source: UCI Machine Learning Repository, <https://archive.ics.uci.edu/ml/datasets/Diabetes+130-US+hospitals>

### VARIABLE : DESCRIPTION

- **Encounter ID** : Unique identifier of an encounter
- **Patient number** : Unique identifier of a patient
- **Race Values** : Caucasian, Asian, African American, Hispanic, and other
- **Gender Values**: male, female, and unknown/invalid
- **Age** : Grouped in 10-year intervals: 0, 10), 10, 20), ..., 90, 100)
- **Weight** : Weight in pounds

- **Admission type** : Integer identifier corresponding to 9 distinct values, for example, emergency, urgent, e
- **Discharge disposition** : Integer identifier corresponding to 29 distinct values, for example, discharged to
- **Admission source** : Integer identifier corresponding to 21 distinct values, for example, physician referral from a hospital
- **Time in hospital** : Integer number of days between admission and discharge
- **Payer code** : Integer identifier corresponding to 23 distinct values, for example, Blue Cross/Blue Shield,
- **Medical specialty** : Integer identifier of a specialty of the admitting physician, corresponding to 84 distinct internal medicine, family/general practice, and surgeon
- **Number of lab procedures** : Number of lab tests performed during the encounter
- **Number of procedures** : Numeric Number of procedures (other than lab tests) performed during the encounter
- **Number of medications** : Number of distinct generic names administered during the encounter
- **Number of outpatient visits** : Number of outpatient visits of the patient in the year preceding the encounter
- **Number of emergency visits** : Number of emergency visits of the patient in the year preceding the encounter
- **Number of inpatient visits** : Number of inpatient visits of the patient in the year preceding the encounter
- **Diagnosis 1** : The primary diagnosis (coded as first three digits of ICD9); 848 distinct values
- **Diagnosis 2** : Secondary diagnosis (coded as first three digits of ICD9); 923 distinct values
- **Diagnosis 3** : Additional secondary diagnosis (coded as first three digits of ICD9); 954 distinct values
- **Number of diagnoses** : Number of diagnoses entered to the system 0%
- **Glucose serum test result** : Indicates the range of the result or if the test was not taken. Values: ">200," measured
- **A1c test result** : Indicates the range of the result or if the test was not taken. Values: ">8" if the result was result was greater than 7% but less than 8%, "normal" if the result was less than 7%, and "none" if not measured
- **Change of medications** : Indicates if there was a change in diabetic medications (either dosage or generic name) "no change"
- **Diabetes medications** : Indicates if there was any diabetic medication prescribed. Values: "yes" and "no"
- 24 features for medications For the generic names: **metformin, repaglinide, nateglinide, chlorpropamide, glipizide, glyburide, tolbutamide, pioglitazone, rosiglitazone, acarbose, miglitol, troglitazone, tolazamide, glyburide-metformin, glipizide-metformin, glimepiride- pioglitazone, metformin-rosiglitazone, and metformin-glipizide** indicates whether the drug was prescribed or there was a change in the dosage. Values: "up" if the dosage increased in the encounter, "down" if the dosage was decreased, "steady" if the dosage did not change, and "no" if the drug was not prescribed
- **Readmitted** : Days to inpatient readmission. Values: "<30" if the patient was readmitted in less than 30 days, "30-90" if readmitted in more than 30 days, and "No" for no record of readmission

## ▼ Overview of Dataset

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

import matplotlib.pyplot as plt
plt.style.use('ggplot')
```

```
# read the file and create a pandas dataframe
```

```
data = pd.read_csv('diabetic_data.csv')
```

```
# check the dimensions of the data  
data.shape
```

```
↳ (101766, 50)
```

```
# take a first look at the dataset  
data.info()
```

```
↳
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101766 entries, 0 to 101765
Data columns (total 50 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   encounter_id                             101766 non-null  int64
1   patient_nbr                             101766 non-null  int64
2   race                                     101766 non-null  object
3   gender                                   101766 non-null  object
4   age                                      101766 non-null  object
5   weight                                  101766 non-null  object
6   admission_type_id                       101766 non-null  int64
7   discharge_disposition_id               101766 non-null  int64
8   admission_source_id                   101766 non-null  int64
9   time_in_hospital                      101766 non-null  int64
10  payer_code                             101766 non-null  object
11  medical_specialty                     101766 non-null  object
12  num_lab_procedures                   101766 non-null  int64
13  num_procedures                       101766 non-null  int64
14  num_medications                      101766 non-null  int64
15  number_outpatient                    101766 non-null  int64
16  number_emergency                     101766 non-null  int64
17  number_inpatient                     101766 non-null  int64
18  diag_1                               101766 non-null  object
19  diag_2                               101766 non-null  object
20  diag_3                               101766 non-null  object
21  number_diagnoses                     101766 non-null  int64
22  max_glu_serum                        101766 non-null  object
23  A1Cresult                            101766 non-null  object
24  metformin                            101766 non-null  object
25  repaglinide                          101766 non-null  object
26  nateglinide                          101766 non-null  object
27  chlorpropamide                       101766 non-null  object
28  glimepiride                           101766 non-null  object
29  acetohexamide                        101766 non-null  object
30  glipizide                             101766 non-null  object
31  glyburide                             101766 non-null  object
32  tolbutamide                           101766 non-null  object
33  pioglitazone                         101766 non-null  object
34  rosiglitazone                        101766 non-null  object
35  acarbose                             101766 non-null  object
36  miglitol                             101766 non-null  object
37  troglitazone                         101766 non-null  object
38  tolazamide                           101766 non-null  object
39  examide                              101766 non-null  object
40  citoglipton                          101766 non-null  object
41  insulin                              101766 non-null  object
42  glyburide-metformin                  101766 non-null  object
43  glipizide-metformin                  101766 non-null  object
44  glimepiride-pioglitazone             101766 non-null  object
45  metformin-rosiglitazone              101766 non-null  object
46  metformin-pioglitazone               101766 non-null  object
47  change                               101766 non-null  object
48  diabetesMed                          101766 non-null  object
49  readmitted                           101766 non-null  object
dtypes: int64(13), object(37)
memory usage: 38.8+ MB

```

```

# first 5 rows of data
data.head()

```




	encounter_id	patient_nbr	race	gender	age	weight	admission_type_id	discharge
0	2278392	8222157	Caucasian	Female	[0-10)	?	6	
1	149190	55629189	Caucasian	Female	[10-20)	?	1	
2	64410	86047875	AfricanAmerican	Female	[20-30)	?	1	
3	500364	82442376	Caucasian	Male	[30-40)	?	1	
4	16680	42519267	Caucasian	Male	[40-50)	?	1	

## ▼ The Response Variable: 'readmitted'

```
# the response variable 'readmitted' in the original dataset contains three categories.  
# 11% of patients were readmitted within 30 days (<30)  
# 35% of patients were readmitted after 30 days (>30)  
# 54% of patients were never readmitted (NO)  
data.groupby('readmitted').size().plot(kind='bar')  
plt.ylabel('Count')
```



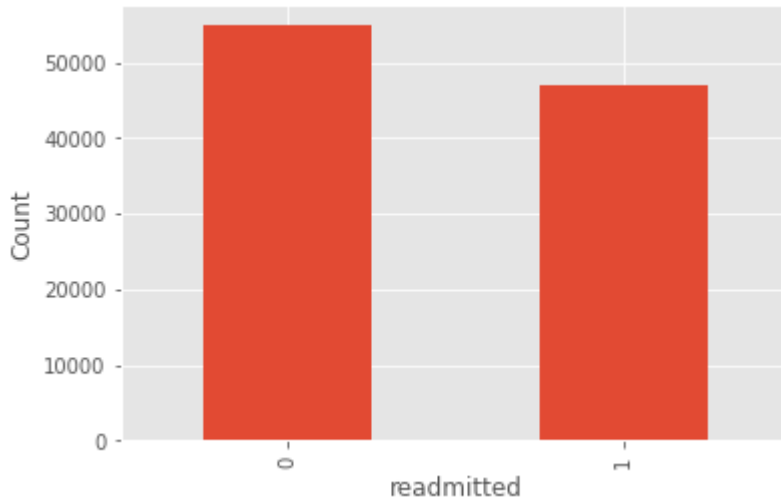
```
Text(0, 0.5, 'Count')
```



```
# recode the response variable so that it becomes a binary classification task.  
# '0' means no readmission  
# '1' means readmission no matter how many days after being discharged  
data['readmitted'] = pd.Series([0 if val == 'NO' else 1 for val in data['readmitted']])  
data_origin = data
```


```
# plot 'readmitted' again to see the change  
data.groupby('readmitted').size().plot(kind='bar')  
plt.ylabel('Count')
```

↗ Text(0, 0.5, 'Count')



```
# the response variable 'discharge_disposition_id' in the original dataset contains 29 categories.  
# now show result  
data.groupby('discharge_disposition_id').size().plot(kind='bar')  
plt.ylabel('Count')
```


↗ Text(0, 0.5, 'Count')



```
# the response variable 'discharge_disposition_id' in the original dataset contains 29 categories.  
# now show result  
data.groupby('change').size().plot(kind='bar')  
plt.ylabel('Count')
```

↗

Text(0, 0.5, 'Count')



## Feature Engineering and EDA

Explore data analysis on each feature and decide how to deal with them.

```
# remove irrelevant features
data.drop(['encounter_id', 'patient_nbr', 'payer_code'], axis=1, inplace=True)
```

```
# check NA in 'weight'
data[data['weight'] == '?'].shape[0] * 1.0 / data.shape[0]
```

```
0.9685847925633315
```

```
# check NA in 'medical_specialty'
data[data['medical_specialty'] == '?'].shape[0] * 1.0 / data.shape[0]
```


```
0.49082208203132677
```

```
# remove 'weight' and 'medical_specialty' because it's hard to do imputation on them
data.drop(['weight', 'medical_specialty'], axis=1, inplace=True)
```

```
# remove rows that have NA in 'race', 'diag_1', 'diag_2', or 'diag_3'
# remove rows that have invalid values in 'gender'
data = data[data['race'] != '?']
data = data[data['diag_1'] != '?']
data = data[data['diag_2'] != '?']
data = data[data['diag_3'] != '?']
data = data[data['gender'] != 'Unknown/Invalid']
```

```
# check 'age' feature
data.groupby('age').size().plot(kind='bar')
plt.ylabel('Count')
```


```
Text(0, 0.5, 'Count')
```



```
# Recategorize 'age' so that the population is more evenly distributed
data['age'] = pd.Series(['[0-50)' if val in ['[0-10)', '[10-20)', '[20-30)', '[30-40)', '[40-50)'] else
                        for val in data['age']], index=data.index)
data['age'] = pd.Series(['[80-100)' if val in ['[80-90)', '[90-100)'] else val
                        for val in data['age']], index=data.index)

data.groupby('age').size().plot(kind='bar')
plt.ylabel('Count')
```

```
➦ Text(0, 0.5, 'Count')
```



```
# original 'discharge_disposition_id' contains 28 levels
# reduce 'discharge_disposition_id' levels into 2 categories
# discharge_disposition_id = 1 corresponds to 'Discharge Home'
data['discharge_disposition_id'] = pd.Series(['Home' if val == 1 else 'Other discharge'
                                              for val in data['discharge_disposition_id']], index=data.index)

# original 'admission_source_id' contains 25 levels
# reduce 'admission_source_id' into 3 categories
data['admission_source_id'] = pd.Series(['Emergency Room' if val == 7 else 'Referral' if val == 1 else 'Other'
                                          for val in data['admission_source_id']], index=data.index)
```



```
# original 'admission_type_id' contains 8 levels
# reduce 'admission_type_id' into 2 categories
data['admission_type_id'] = pd.Series(['Emergency' if val == 1 else 'Other type'
                                       for val in data['admission_type_id']], index=data.index)
```

23 features represent different medications, should we use them all in the model?

How to deal with 'diag\_1', 'diag\_2', and 'diag\_3'?

Based on the results found in Reference[1], we know that:

- The statistical model suggests that the relationship between the probability of readmission and the HbA1c is **primary diagnosis**.
- The only medication that varied significantly across the patients was the delivery of **insulin**, while other medications were not significantly different among all the patients.
- It may not be surprising that the attention given to diabetes care in individuals with admitting diagnoses of diabetes mellitus may have been less than those with a **primary diagnosis of diabetes mellitus**.
- Greater attention to diabetes care during the hospitalization for these high-risk individuals may have a significant impact on readmission rates.
- The primary, secondary, and third medical diagnoses are marked by the ICD9 codes.
- ICD9 code for diabetes: 250.xx

```
# compare diabetes medications 'miglitol', 'nateglinide' and 'acarbose' with 'insulin', as an example
fig = plt.figure(figsize=(20,15))

ax1 = fig.add_subplot(221)
ax1 = data.groupby('miglitol').size().plot(kind='bar')
plt.xlabel('miglitol', fontsize=15)
plt.ylabel('Count', fontsize=15)

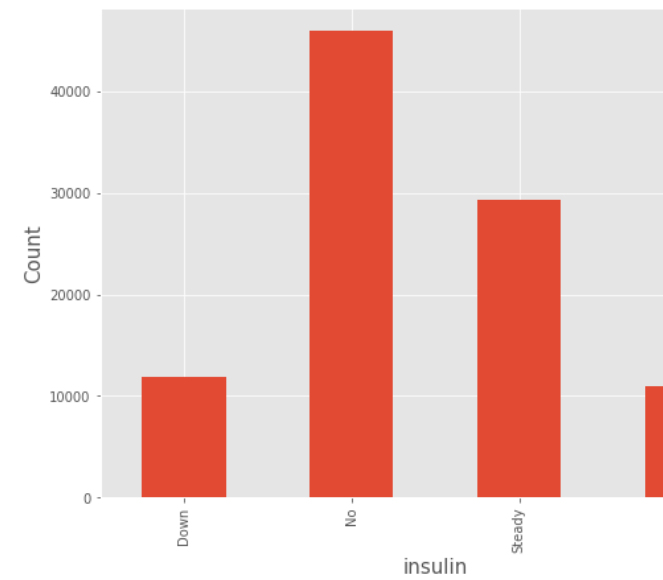
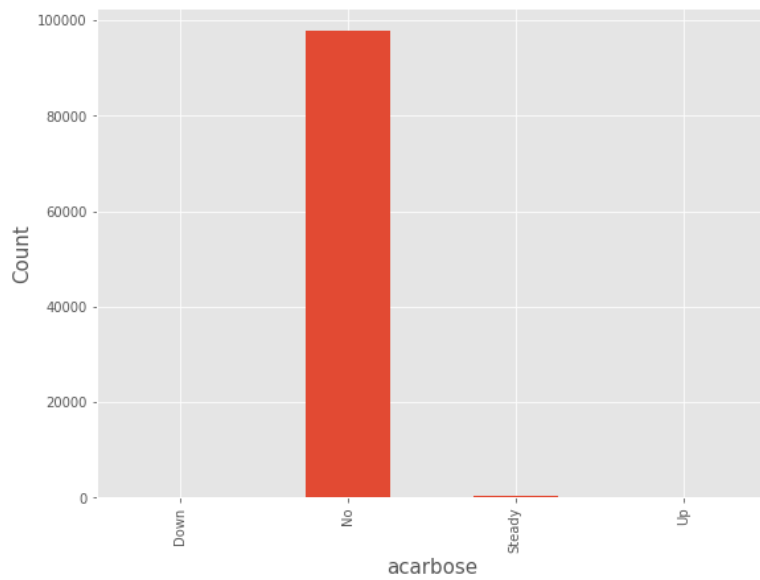
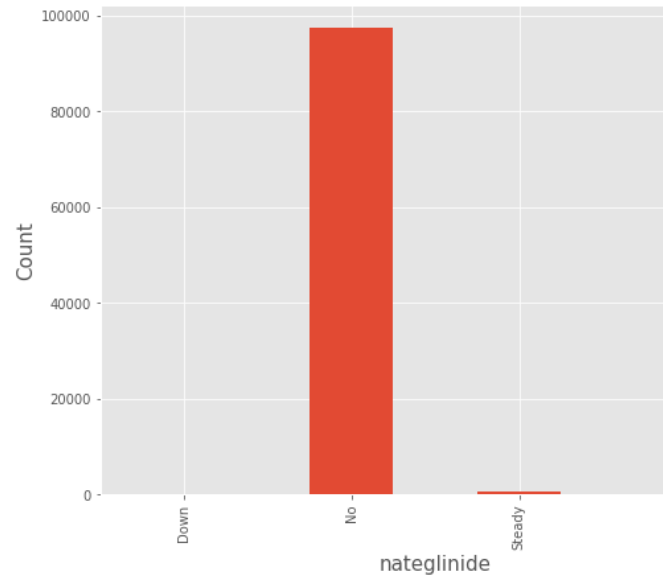
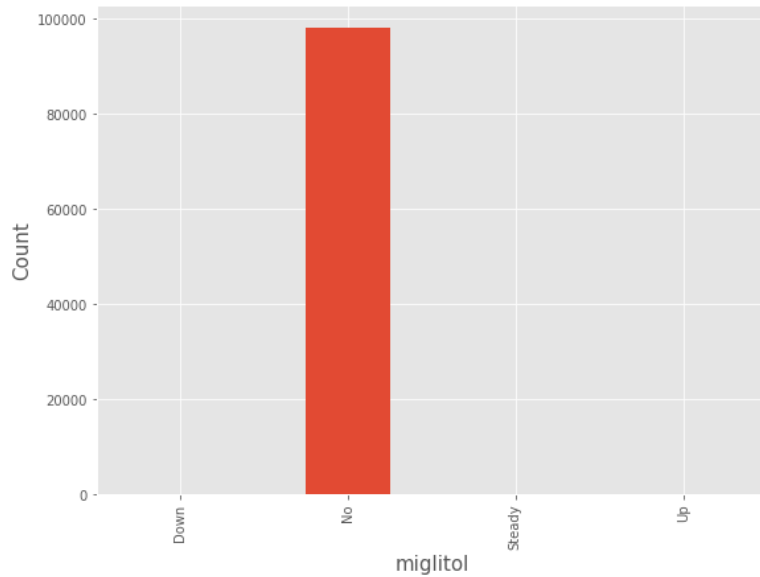
ax2 = fig.add_subplot(222)
ax2 = data.groupby('nateglinide').size().plot(kind='bar')
plt.xlabel('nateglinide', fontsize=15)
plt.ylabel('Count', fontsize=15)

ax3 = fig.add_subplot(223)
ax3 = data.groupby('acarbose').size().plot(kind='bar')
plt.xlabel('acarbose', fontsize=15)
plt.ylabel('Count', fontsize=15)

ax4 = fig.add_subplot(224)
ax4 = data.groupby('insulin').size().plot(kind='bar')
plt.xlabel('insulin', fontsize=15)
plt.ylabel('Count', fontsize=15)
```



Text(0, 0.5, 'Count')



```
# keep only 'insulin' and remove the other 22 diabetes medications
data.drop(['metformin', 'repaglinide', 'nateglinide', 'chlorpropamide', 'glimepiride',
          'acetoexamide', 'glipizide', 'glyburide', 'tolbutamide', 'pioglitazone',
          'rosiglitazone', 'acarbose', 'miglitol', 'troglitazone', 'tolazamide', 'examide',
          'citoglipton', 'glyburide-metformin', 'glipizide-metformin', 'glimepiride-pioglitazone',
          'metformin-rosiglitazone', 'metformin-pioglitazone'], axis=1, inplace=True)
```

```
# denote 'diag_1' as '1' if it relates to diabetes and '0' if it's not
# remove 'diag_2' and 'diag_3'
data['diag_1'] = pd.Series([1 if val.startswith('250') else 0 for val in data['diag_1']], index=data.index)
data.drop(['diag_2', 'diag_3'], axis=1, inplace=True)
```

```
data.shape
```

```
↳ (98052, 21)
```

```
# we end up with 98052 rows, 20 features, and one response variable 'readmitted'
list(data.columns.values)
```

```
↳ ['race',
    'gender',
    'age',
    'admission_type_id',
    'discharge_disposition_id',
    'admission_source_id',
    'time_in_hospital',
    'num_lab_procedures',
    'num_procedures',
    'num_medications',
    'number_outpatient',
    'number_emergency',
    'number_inpatient',
    'diag_1',
    'number_diagnoses',
    'max_glu_serum',
    'A1Cresult',
    'insulin',
    'change',
    'diabetesMed',
    'readmitted']
```

## ▼ Data Preprocessing

```
# one-hot-encoding on categorical features
# convert nominal values to dummy values
# make age group
data['age'] = pd.Series([25 if val == '[0-50)' else 55 if val == '[50-60)' else 65 if val == '[60-70)' else 75 for val in data['age']], index=data.index)
```

```
df_race = pd.get_dummies(data['race'])
df_gender = pd.get_dummies(data['gender'])
df_max_glu_serum = pd.get_dummies(data['max_glu_serum'])
df_A1Cresult = pd.get_dummies(data['A1Cresult'])
df_insulin = pd.get_dummies(data['insulin'])
df_change = pd.get_dummies(data['change'])
df_diabetesMed = pd.get_dummies(data['diabetesMed'])
df_discharge_disposition_id = pd.get_dummies(data['discharge_disposition_id'])
df_admission_source_id = pd.get_dummies(data['admission_source_id'])
df_admission_type_id = pd.get_dummies(data['admission_type_id'])
```

```
data = pd.concat([data, df_race, df_gender, df_max_glu_serum, df_A1Cresult,
                  df_insulin, df_change, df_diabetesMed, df_discharge_disposition_id,
```

```
df_admission_source_id, df_admission_type_id], axis=1)
data.drop(['race', 'gender', 'max_glu_serum', 'A1Cresult', 'insulin', 'change',
          'diabetesMed', 'discharge_disposition_id', 'admission_source_id',
          'admission_type_id'], axis=1, inplace=True)
```

```
data.head()
```



	age	time_in_hospital	num_lab_procedures	num_procedures	num_medications	number_outpatient
1	25	3	59	0	18	
2	25	2	11	5	13	
3	25	2	44	1	16	
4	25	1	51	0	8	
5	55	3	31	6	16	

```
# feature scaling, features are standardized to have zero mean and unit variance
feature_scale_cols = ['time_in_hospital', 'num_lab_procedures', 'num_procedures', 'num_medications',
                     'number_diagnoses', 'number_inpatient', 'number_emergency', 'number_outpatient']
```

```
from sklearn import preprocessing
scaler = preprocessing.StandardScaler().fit(data[feature_scale_cols])
data_scaler = scaler.transform(data[feature_scale_cols])
```

```
data_scaler_df = pd.DataFrame(data=data_scaler, columns=feature_scale_cols, index=data.index)
data.drop(feature_scale_cols, axis=1, inplace=True)
data = pd.concat([data, data_scaler_df], axis=1)
```

## Data Preparation Summary

1. Remove NA in 'race', 'gender', 'diag\_1', 'diag\_2', and 'diag\_3'
2. Remove 29 features in total: 'encounter\_id', 'patient\_nbr', 'payer\_code', 'weight', 'medical\_specialty', 'diag\_2', 'diag\_3', and 'metformin', 'repaglinide', 'nateglinide', 'chlorpropamide', 'glimepiride', 'acetohexamide', 'glipizide', 'glyburide', 'tolbutamide', 'acarbose', 'miglitol', 'troglitazone', 'tolazamide', 'examide', 'citoglipton', 'glyburide-metformin', 'glipizide-metformin', 'glim', 'rosiglitazone', and 'metformin-pioglitazone'
3. Recategorize 'age' feature
4. Reduce levels in 'discharge\_disposition\_id', 'admission\_source\_id', and 'admission\_type\_id'
5. One-hot-encode on categorical data
6. Square root transform on right skewed count data
7. Apply feature standardizing on numerical data

## Feature Importance

Evaluate the importance of different features by using ensemble method.

```
# create X (features) and y (response)
X = data.drop(['readmitted'], axis=1)
```

```
y = data['readmitted']
```

```
# split X and y into cross-validation (85%) and testing (15%) data sets
from sklearn.model_selection import train_test_split
X_cv, X_test, y_cv, y_test = train_test_split(X, y, test_size=0.15)
```

```
# fit Random Forest model to the cross-validation data
from sklearn.ensemble import RandomForestClassifier
forest = RandomForestClassifier()
forest.fit(X_cv, y_cv)
importances = forest.feature_importances_
```

```
# make importance relative to the max importance
feature_importance = 100.0 * (importances / importances.max())
sorted_idx = np.argsort(feature_importance)
feature_names = list(X_cv.columns.values)
feature_names_sort = [feature_names[indice] for indice in sorted_idx]
pos = np.arange(sorted_idx.shape[0]) + .5
print ('Top 10 features are: ')
for feature in feature_names_sort[::-1][:10]:
    print (feature)
```


```
# plot the result
plt.figure(figsize=(12, 10))
plt.barh(pos, feature_importance[sorted_idx], align='center')
plt.yticks(pos, feature_names_sort)
plt.title('Relative Feature Importance', fontsize=20)
plt.show()
```



Top 10 features are:

num\_lab\_procedures  
num\_medications  
time\_in\_hospital  
age  
number\_diagnoses  
num\_procedures  
number\_inpatient  
number\_outpatient  
number\_emergency  
Female

Relative Feature Importance



```
# make a smaller feature set which only contains the top 10 features
X_cv_top10 = X_cv[['time_in_hospital', 'num_lab_procedures', 'num_procedures', 'num_medications', 'number_inpatient', 'number_outpatient', 'number_emergency', 'Female', 'age']]
X_test_top10 = X_test[['time_in_hospital', 'num_lab_procedures', 'num_procedures', 'num_medications', 'number_inpatient', 'number_outpatient', 'number_emergency', 'Female', 'age']]
data.dtypes
```



age	int64
diag_1	int64
readmitted	int64
AfricanAmerican	uint8
Asian	uint8
Caucasian	uint8
Hispanic	uint8
Other	uint8
Female	uint8
Male	uint8
>200	uint8
>300	uint8
None	uint8
Norm	uint8
>7	uint8
>8	uint8
None	uint8
Norm	uint8
Down	uint8
No	uint8
Steady	uint8
Up	uint8
Ch	uint8
No	uint8
No	uint8
Yes	uint8
Home	uint8
Other discharge	uint8
Emergency Room	uint8
Other source	uint8
Referral	uint8
Emergency	uint8
Other type	uint8
time_in_hospital	float64
num_lab_procedures	float64
num_procedures	float64
num_medications	float64
number_diagnoses	float64
number_inpatient	float64
number_emergency	float64
number_outpatient	float64
dtype:	object

## ▼ Model Selection

Compare the performance of three machine learning algorithms: Random Forest, Naive Bayes, and Logistic Regression.

Select the best model using 10-fold cross validation. The metric that will be using is the accuracy of prediction.

```
from sklearn.model_selection import cross_val_score
```

```
# Random Forest
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
clf1 = RandomForestClassifier()
```

```
RF_score = cross_val_score(clf1, X_cv_top10, y_cv, cv=10, scoring='accuracy').mean()
```

```
RF_score
```

0.5893984075709607

```
# Naive Bayes
from sklearn.naive_bayes import GaussianNB
clf2 = GaussianNB()
NB_score = cross_val_score(clf2, X_cv_top10, y_cv, cv=10, scoring='accuracy').mean()
NB_score
```

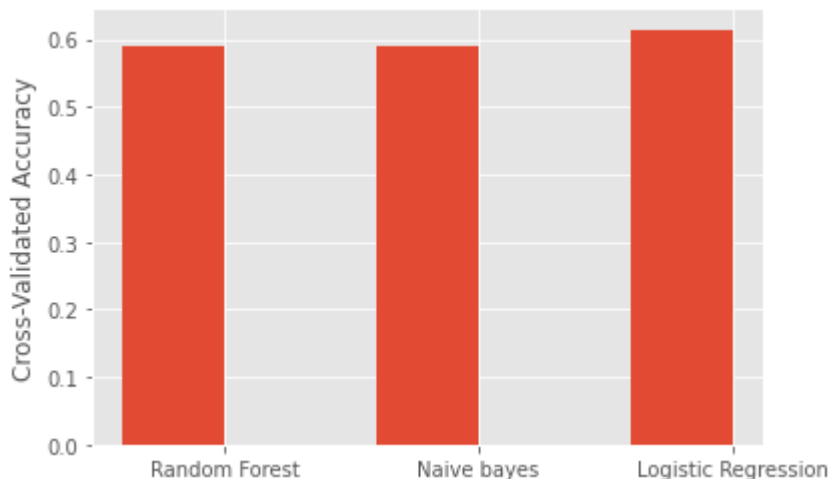
0.5892709345949001

```
# Logistic Regression
from sklearn.linear_model import LogisticRegression
clf3 = LogisticRegression()
LR_score = cross_val_score(clf3, X_cv_top10, y_cv, cv=10, scoring='accuracy').mean()
LR_score
```

0.614410889658734

```
# plot and compare the scores
# LR outperforms the other two a little bit
x_axis = np.arange(3)
y_axis = [RF_score, NB_score, LR_score]
plt.bar(x_axis, y_axis, width=0.4)
plt.xticks(x_axis + 0.4/2., ('Random Forest', 'Naive bayes', 'Logistic Regression'))
plt.ylabel('Cross-Validated Accuracy')
```

Text(0, 0.5, 'Cross-Validated Accuracy')



```
# DecisionTree
from sklearn.tree import DecisionTreeClassifier
clf4 = DecisionTreeClassifier(random_state=0, max_depth=32, criterion = "entropy", min_samples_split=
De_score = cross_val_score(clf4, X_cv_top10, y_cv, cv=10, scoring='accuracy').mean()
De_score
```

0.553728205098204

```
# k means
from sklearn.neighbors import KNeighborsClassifier
clf5 = KNeighborsClassifier(12)
Kn_score = cross_val_score(clf5, X_cv_top10, y_cv, cv=10, scoring='accuracy').mean()
Kn_score
```



0.5888757784444858

```
# neural network keras
```

```
import keras
```

```
from tensorflow.python.keras.layers import Dense
```

```
from tensorflow.python.keras import Sequential
```

```
classifier = Sequential()
```

```
classifier.add(Dense(input_shape=(10,), units=6, kernel_initializer='random_uniform', activation='re
```

```
classifier.add(Dense(units=10, kernel_initializer='random_uniform', activation='relu'))
```

```
classifier.add(Dense(units=15, kernel_initializer='random_uniform', activation='relu'))
```

```
classifier.add(Dense(units=1, kernel_initializer='random_uniform', activation='sigmoid'))
```

```
classifier.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
classifier.fit(X_cv_top10, y_cv, batch_size=10, epochs=3)
```

```
y_pred_nn = classifier.predict(X_test_top10)
```

```
y_pred_nn = [1 if y > 0.5 else 0 for y in y_pred_nn]
```

```
from sklearn.metrics import accuracy_score
```

```
acc1 = accuracy_score(y_test, y_pred_nn)
```

```
acc1
```

Epoch 1/3

8335/8335 [=====] - 12s 1ms/step - loss: 0.6608 - accuracy: 0.6061

Epoch 2/3

8335/8335 [=====] - 11s 1ms/step - loss: 0.6563 - accuracy: 0.6125

Epoch 3/3

8335/8335 [=====] - 11s 1ms/step - loss: 0.6554 - accuracy: 0.6153

0.6231982594506391

other solution for classtering this dataset

<https://www.kaggle.com/iabhishekofficial/prediction-on-hospital-readmission/notebook> link

